

23.1 More on B-trees

23.1.1 Height of a B-tree

The height of a tree with n nodes is $\Theta((\log n)/(\log M))$. For the exam, know the algorithms that can be applied on this tree (delete, search, insert), and also know that the height of a tree is $\Theta((\log n)/(\log M))$.

Proof of the height of a B-tree

M is the maximum number of children for any node. With that in mind, observe that the number of keys n in the tree is

$$\begin{aligned} n &\geq \left(\frac{M}{2}\right)^h \\ \log(n) &\geq h \log\left(\frac{M}{2}\right) \\ \frac{\log(n)}{\log(\frac{M}{2})} &\geq h \\ h &\in O\left(\frac{\log(n)}{\log(M)}\right) \end{aligned}$$

Similarly we have:

$$\begin{aligned} n &\leq M^{h+1} \\ \log(n) &\leq (h+1) \log(M) \\ h &\in \Omega\left(\frac{\log(n)}{\log(M)}\right) \end{aligned}$$

Thus, we have $h \in \Theta\left(\frac{\log(n)}{\log(M)}\right)$

23.1.2 Hashing in External Memory

How can we hash and minimize disk transfers? We can use **extendible hashing**, which is similar to a B-tree with height 1 and max size S at the leaves.

Structure

The directory is stored in internal memory. It contains an array of size 2^d , where d is called the *order*. Each directory entry points to a block stored in external memory. Each block contains at most S items. Note that many entries can point to the same block. The values in every block in block B agree on leading k_B bits: Know how the data structure is constructed. How to search. How to insert items. FORGET ABOUT DELETE.

